

Docker Basics

Agenda

- What is containerization?
- Why we need containerization?
- What is docker?
- Installing Docker (Linux, Mac only)
- Dockerfile Jupyter example - creating and running

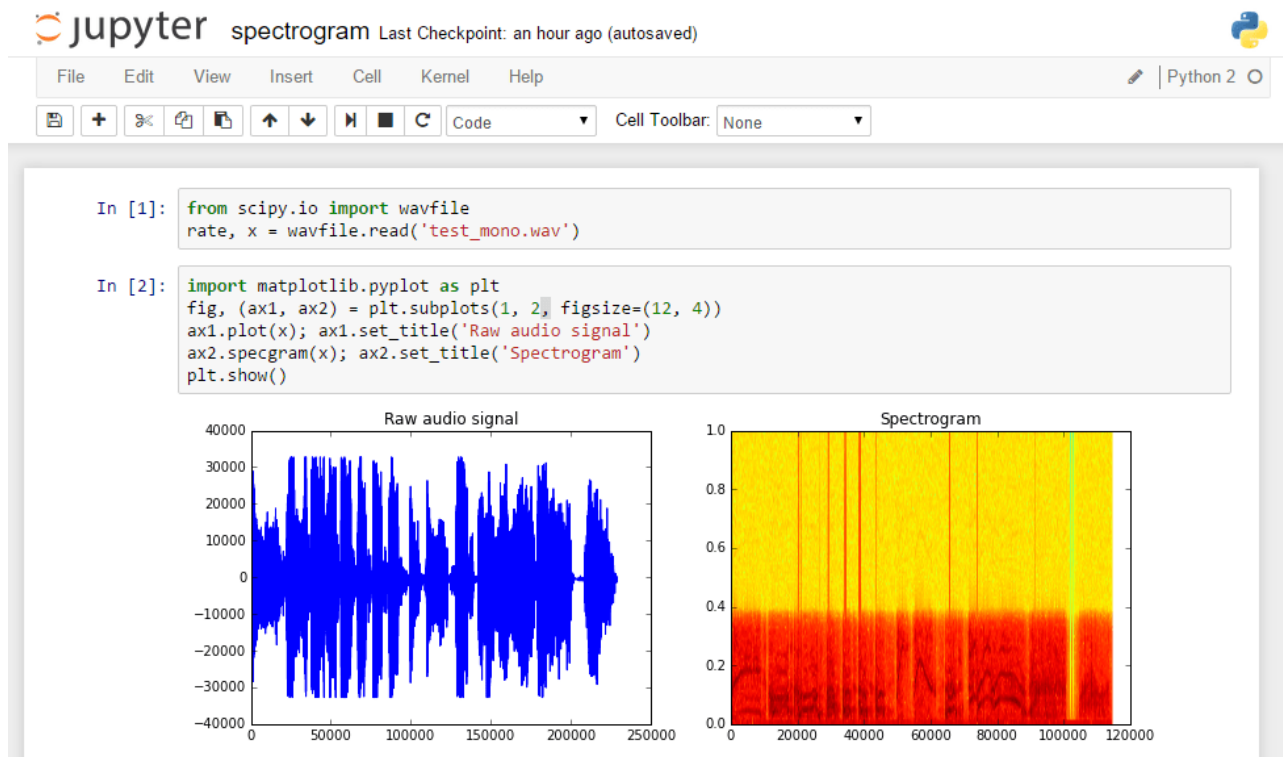
What is containerization ?

- Containers are programs with all dependencies installed. (like those of windows, where you click Next)
- Like Virtual-Machines, but lighter.



Why would you need containers?

- Analysis with Jupyter notebook which you want to share with your friends, so they can tinker around with it and run it themselves.



Issues

- First they will have to setup/install correct libraries (specific versions)
- Everyone might be on different OS - different instructions for all (Linux, Windows, Mac)
- Each person had to do that by himself

Containers Solve these Issues

- Software Dependency Resolution (OS, Libraries)
- No Setup Time

What is Docker?

- Company which provides this containerization technology - a synonym for containerization



- Other compaies do exist :
 - rkt (aka rocket)
 - Mesos Containers
 - Windows Server Containers

Installing Docker

Ubuntu

```
In [ ]: sudo apt-get update
# install
sudo apt install docker.io
# enable docker when system boots up
sudo systemctl start docker
sudo systemctl enable docker
# check version
docker --version
```

MAC

Go to this link and download the stable version

```
In [ ]: https://docs.docker.com/v17.12/docker-for-mac/install/#download-docker-for-mac
```

Demo Create containers

In []: docker pull ubuntu

```
sahil@sahil-Inspiron-5570:~$ docker pull ubuntu
Using default tag: latest
latest: Pulling from library/ubuntu
5b7339215d1d: Pull complete
14ca88e9f672: Pull complete
a31c3b1caad4: Pull complete
b054a26005b7: Pull complete
Digest: sha256:9b1702dcfe32c873a770a32cfd306dd7fc1c4fd134adfb783db68defc8894b3c
Status: Downloaded newer image for ubuntu:latest
sahil@sahil-Inspiron-5570:~$
```


In []: docker images

```
sahil@sahil-Inspiron-5570:~$ docker images
REPOSITORY              TAG
nginx                    latest
dockeriseddataprocessingpipeline_consumer  latest
dockeriseddataprocessingpipeline_producer   latest
wurstmeister/kafka       2.12-2.2.1
postgres                  latest
ubuntu                    latest
python                    3.7-slim
das_notebook              latest
```

TAG	IMAGE ID	CREATED	SIZE
latest	f68d6e55e065	10 days ago	109MB
latest	9edb9b8bd634	12 days ago	329MB
latest	9edb9b8bd634	12 days ago	329MB
2.12-2.2.1	0a993d8130df	13 days ago	421MB
latest	79db2bf18b4a	2 weeks ago	312MB
latest	4c108a37151f	3 weeks ago	64.2MB
3.7-slim	338ae06dfca5	4 weeks ago	143MB
latest	22c3e2374585	4 weeks ago	4.72GB

In []: docker run -it ubuntu bash

```
sahil@sahil-Inspiron-5570:~$ docker run -it ubuntu bash
root@94f45ba5ad87:/# ls
bin boot dev etc home lib lib64 media mnt opt proc root run sbin srv sys  usr var
root@94f45ba5ad87:/# pwd
/
root@94f45ba5ad87:/# cd home/
root@94f45ba5ad87:/home# ls
root@94f45ba5ad87:/home#
```

In []: docker ps

```
sahil@sahil-Inspiron-5570:~$ docker run -it ubuntu bash
root@94f45ba5ad87:/# ls
bin boot dev etc home lib lib64 media mnt opt proc root run sbin srv sys usr var
root@94f45ba5ad87:/# pwd
/
root@94f45ba5ad87:/# cd home/
root@94f45ba5ad87:/home# ls
root@94f45ba5ad87:/home#
```

```
Terminal
sahil@sahil-Inspiron-5570:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
94f45ba5ad87	ubuntu	"bash"	About a minute ago	Up About a minute		silly_leavitt
9f547b72c645	das_notebook	"jupyter lab --ip=0.0.0.0"	6 minutes ago	Up 6 minutes	0.0.0.0:4040->4040/tcp, 0.0.0.0:8888->8888/tcp	flanboyant_dhawan
74e74d0b10aa	587aa1d0e586	"docker-entrypoint.s..."	6 weeks ago	Restarting (1) 43 seconds ago		postgres_postgres_1
b6c733d904f4	redis	"docker-entrypoint.s..."	2 months ago	Up 10 minutes	6379/tcp	postgres_redis_1

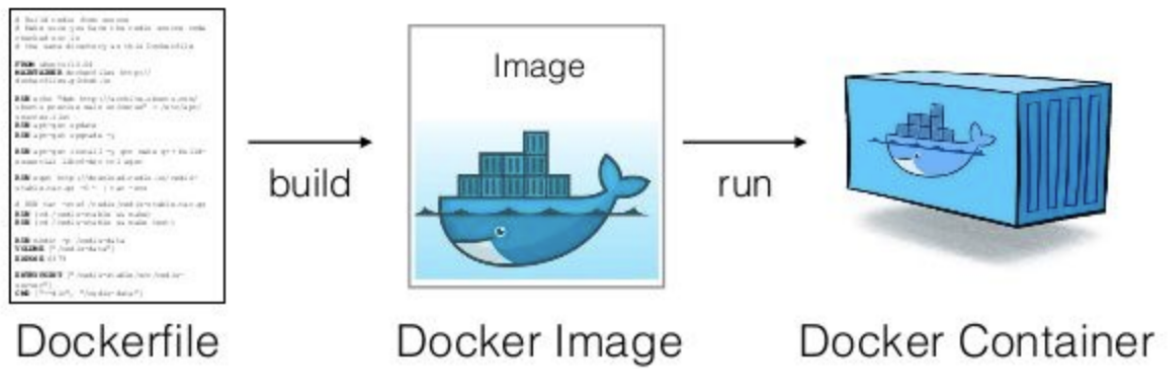
```
sahil@sahil-Inspiron-5570:~$
```

Try out different docker images

- `docker pull python`
- `docker pull postgres`
- `docker pull`

Dockerfiles

Dockerfile (text file) create docker images which we run as containers



Lets Create our Own Jupyter Dockerfile

Create a file called **Dockerfile** and paste the following lines,

```
In [ ]: FROM ubuntu

        RUN apt-get update
        RUN apt-get upgrade -y

        RUN apt-get install -y python3 python3-dev python3-pip
```

Create the image using,

```
In [ ]: docker build . -t my_python_image
```

Check if the image is built using,

```
In [ ]: docker images
```

Now we can install all our popular python packages

Rebuild the image using the same command as previous time

```
In [ ]: FROM ubuntu  
  
RUN apt-get update  
RUN apt-get upgrade -y  
  
RUN apt-get install -y python3 python3-dev python3-pip  
  
RUN pip3 install jupyter pandas numpy scipy
```

You will notice that only the new line is run. All previous instructions are loaded from cache. So only new changes are run henceforth.

```
sahil@sahil-Inspiron-5570:~/../docker-basics-workshop$ docker build . -t my_python_image
Sending build context to Docker daemon 697.3kB
Step 1/5 : FROM ubuntu
--> 4c108a37151f
Step 2/5 : RUN apt-get update
--> Using cache
--> 6e5b2d780d3e
Step 3/5 : RUN apt-get upgrade -y
--> Using cache
--> 9673e2b637e5
Step 4/5 : RUN apt-get install -y python3 python3-dev python3-pip
--> Running in 67b8ac322b02
Reading package lists...
Building dependency tree...
Reading state information...
The following additional packages will be installed:
```


Lets run the Docker container with jupyter notebook

Add the entrypoint command and rebuild the image

```
In [ ]: FROM ubuntu

RUN apt-get update
RUN apt-get upgrade -y

RUN apt-get install -y python3 python3-dev python3-pip

RUN pip3 install jupyter numpy

ENTRYPOINT ["jupyter", "notebook", "--ip=0.0.0.0", "--allow-root", "--port=8889"]
```

Run the container

We have mapped the port 8889 which is on the container to port 8889 on the host

```
In [ ]: docker run -p HOST_PORT:CONTAINER_PORT image-name
```

```
In [ ]: docker run -p 8889:8889 my_python_image
```

```
sahil@sahil-Inspiron-5570:~/.../docker-basics-workshop$ docker run -p 8889:8889 my_python_image
[I 12:31:15.991 NotebookApp] Writing notebook server cookie secret to /root/.local/share/jupyter/runtime/notebook_cookie_secret
[I 12:31:16.152 NotebookApp] Serving notebooks from local directory: /
[I 12:31:16.152 NotebookApp] The Jupyter Notebook is running at:
[I 12:31:16.152 NotebookApp] http://(163aabd1d8e6 or 127.0.0.1):8889/?token=8b8537b390a311a660ecb8d05d61d6103c65e66d306a3e39
[I 12:31:16.152 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[W 12:31:16.156 NotebookApp] No web browser found: could not locate runnable browser.
[C 12:31:16.156 NotebookApp]






To access the notebook, open this file in a browser:
    file:///root/.local/share/jupyter/runtime/nbserver-1-open.html
Or copy and paste one of these URLs:
    http://(163aabd1d8e6 or 127.0.0.1):8889/?token=8b8537b390a311a660ecb8d05d61d6103c65e66d306a3e39
^[[24~
```

Persisting Data

If we want to save our files, we will attach a folder from our computer to another folder inside the docker container. Create a directory called **my-data** in your computer, we will store all our jupyter notebook analysis in here.

```
In [ ]: docker run -v /HOST/DIRECTORY/FULL/PATH:/container/directory image-name
```

```
In [ ]: docker run -p 8889:8889 -v "$(pwd)/my-data":"/home/ my_python_image
```

Name	
	images
	my-data
	Dockerfile
	presentation.ipynb
	presentation.slides.html